

process

Chapter 4: Multithreaded Programming

Prof. Li-Pin Chang
National Chiao Tung University

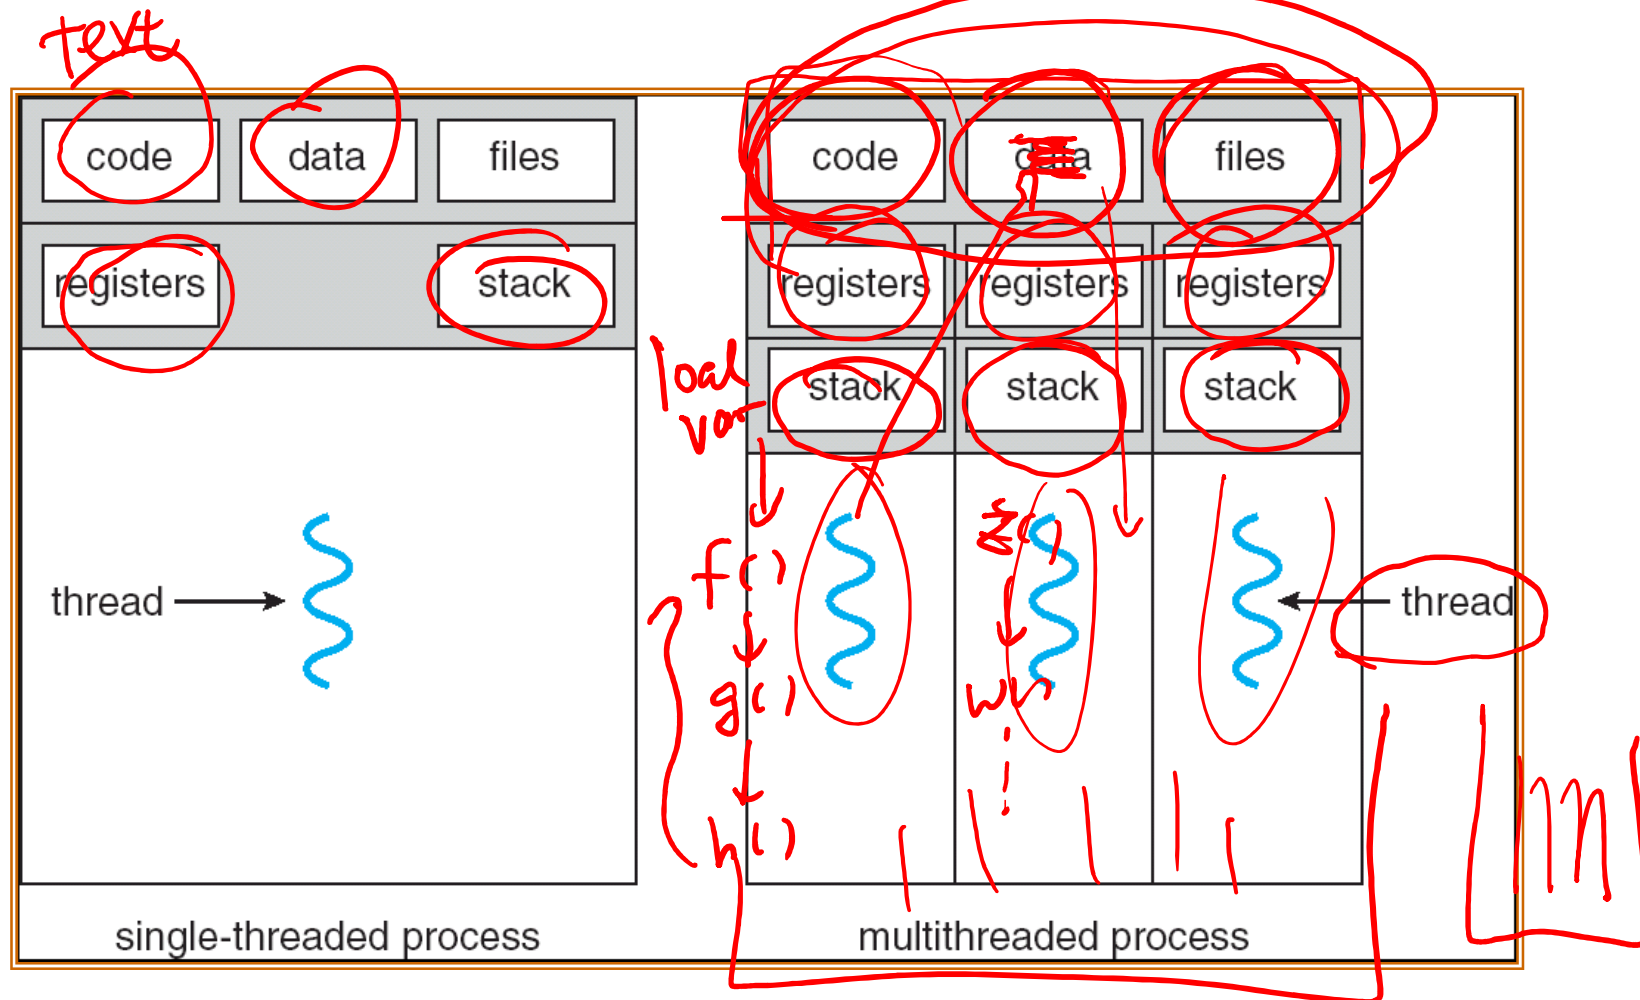
Chapter 4: Multithreaded Programming

- Overview
- Multithreading Models
- Thread Libraries
- Threading Issues
- Operating-System Examples

3rd


OVERVIEW

Single and Multithreaded Processes



- A process is a “container” of all its threads

Benefits



- Responsiveness

- A thread accepts UI inputs while another does computation

- Resource Sharing

- To share code and most of the data structures

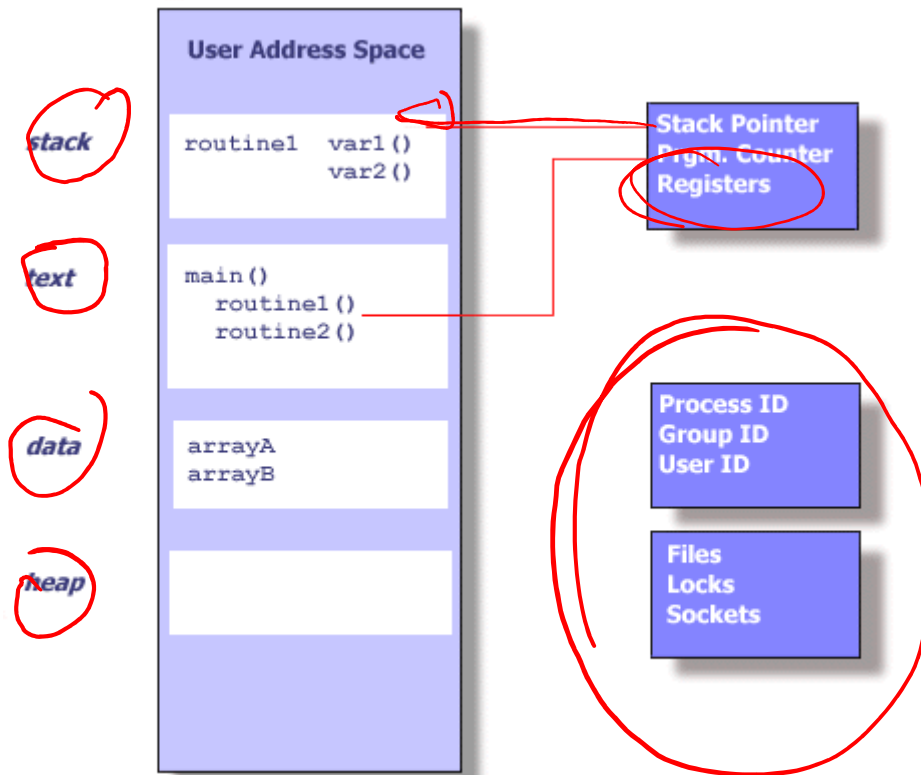
- Economy

- A thread is a lightweight process

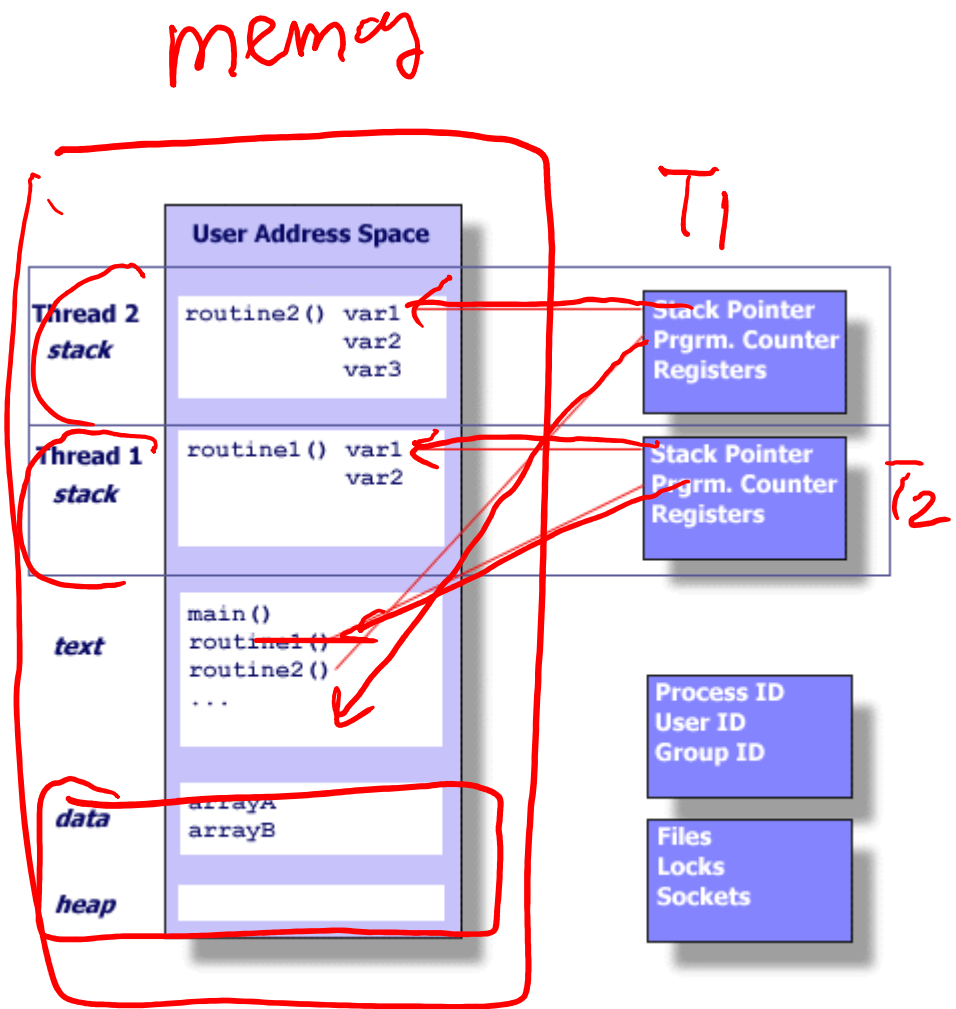
- Utilization of MP Architectures

- To utilize multiple cores or to improve ILP

In Solaris it is 5 times slower to context switch a process than to context switch a thread, and 13 times slower for creation.



A process



Two threads in a process

- This independent flow of control is accomplished because a thread maintains its own:
 - Stack pointer
 - Registers
 - Scheduling properties (such as policy or priority)
 - Set of pending and blocked signals
 - Thread specific data.

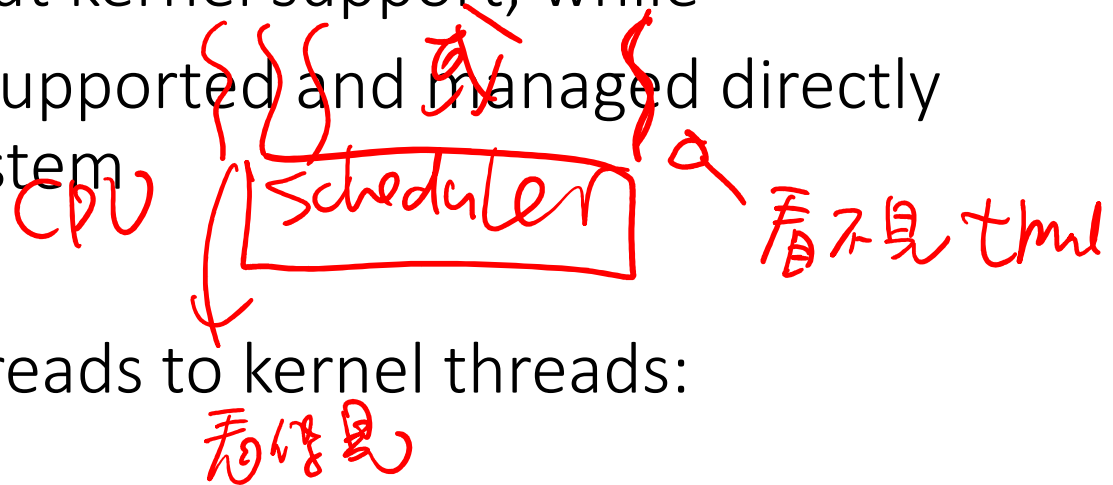

MULTITHREADING MODELS

Multithreading Models

- User threads are supported above the kernel and are managed without kernel support, while
- kernel threads are supported and managed directly by the operating system

- Mapping of user threads to kernel threads:

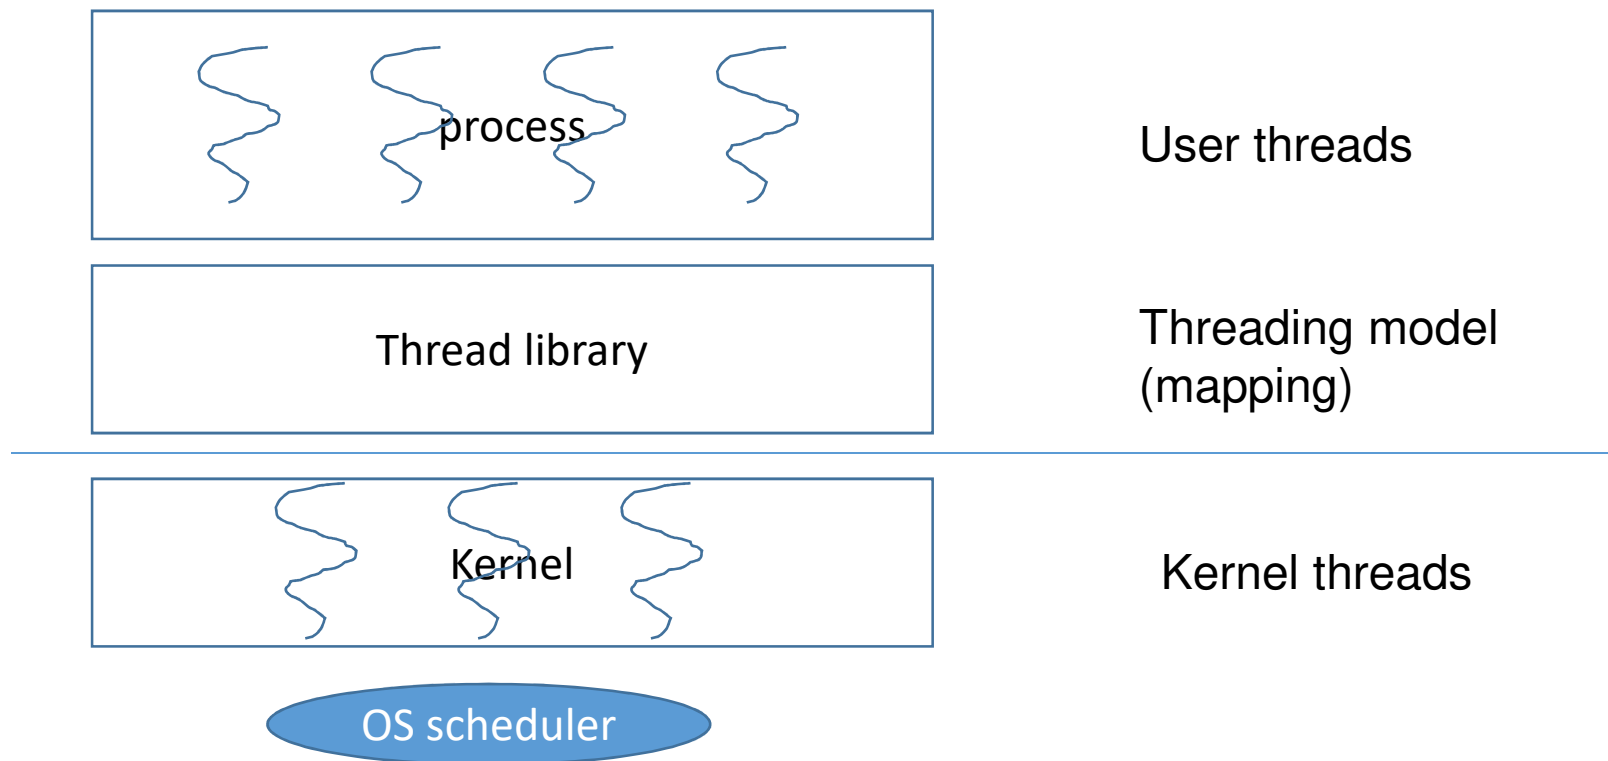
- Many-to-One
- One-to-One
- Many-to-Many



看不見 thread

看得見

Multithreading Models



Specification vs. Implementation

- Major thread libraries
 - Pthreads
 - Win32 threads
 - Java threads
- Pthread is a specification (part of POSIX)
 - How Pthread is implemented (which threading model) is not part of the Pthread specification
 - Pthead can be implemented using 1-1, M-1, or M-M

Many-to-One

- Many user-level threads mapped to single kernel thread

- Examples:

- Solaris Green Threads


- For JDK. Options: green or native

- GNU Portable Threads

- An implementation of the Pthread specification

~~M:1~~
Pthread → spec

Many-to-One Model



- If one single thread issues a blocking system call, then the entire collection of threads become blocked
- Not MP-aware
- Not truly concurrent

Many-to-One Model

- Use special wrapper functions to prevent blocking (sync) calls from stalling all user threads

For example, use

`ssize_t pth_read(int fd, void *buf, size_t nbytes);`

Instead of `read()` *blocking call*


For `GNU portable threads (M-1)` only

T Lib

One-to-One

- Each user-level thread maps to kernel thread
- Examples
 - Windows NT/XP/2000
 - Linux Pthread
 - Solaris 9 and later

One-to-one Model



glibc


kernel dependency

- Higher management overheads for threads
- MP-aware
- A blocking call will not block the whole thread set

Many-to-Many Model

- Allows many user level threads to be mapped to many kernel threads
 - One thread won't block the entire process
- Allows the operating system to create a sufficient number of kernel threads
 - More economic than 1-1 model
- Solaris prior to version 9
- Windows NT/2000 with the ThreadFiber package

Many-to-Many Model




- A mixture of 1-1 and M-1

Two-level Model

- Similar to M:M, except that static binding between user threads and kernel threads is permitted
- Examples
 - IRIX
 - HP-UX
 - Tru64 UNIX
 - Solaris 8 and earlier

Two-level Model




Threading Models

- Specifications do not cover threading model
 - Pthread may be many-to-one, one-to-one, or many-to-many
 - Java green threads are many-to-one but Java threads are not defined
- Different thread library may adopt different threading models
 - GNU Portable Thread is many-to-one
 - Linux Pthread?
- You should not assume the threading model of a thread library
 - Check the programmers' manual first, or write a small test program to make sure

Light-Weight Processes

- LWP is an **optional** abstraction of kernel scheduling units
- An LWP is like a virtual processor on which user threads are scheduled
- Basically the mapping of LWPs to kernel threads is 1-1
- The mapping of user threads to LWP is 1-1, M-1, or M-M



- Consider a multithreaded process consisting of 2 IO-bound threads and 4 CPU-bound threads. Let the threading model be M:M. Let there be 5 kernel threads.

At most how many CPU cores that the multithreaded process can fully utilize?

- What if there were 2 I/O-bound threads and 4 CPU-bound threads?

THREAD LIBRARIES

Pthread

- =POSIX thread
- Pthread is a “specification”, not an implementation
- Implementations:
 - GNU portable thread (M-1)
 - Linux Pthread (1-1)
 - Mac OS X Pthread (?)

```
#include <pthread.h>
#include <stdio.h>
```

-pthread

1100

```
int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* the thread */
```

```
int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    if (argc != 2) {
        fprintf(stderr, "usage: a.out <integer value>\n");
        return -1;
    }
    if (atoi(argv[1]) < 0) {
        fprintf(stderr, "%d must be >= 0\n", atoi(argv[1]));
        return -1;
    }

```

```
/* get the default attributes */
pthread_attr_t attr;
pthread_attr_init(&attr);
```

```
/* create the thread */
```

```
pthread_create(&tid, &attr, runner, argv[1]);
```

```
/* wait for the thread to exit */
```

```
pthread_join(tid, NULL);
```

```
printf("sum = %d\n", sum);
```

```
/* The thread will begin control in this function */
```

```
void *runner(void *param)
```

```
{
```

```
    int i, upper = atoi(param);
```

```
    sum = 0;
```

```
    for (i = 1; i <= upper; i++)
```

```
        sum += i;
```

```
    pthread_exit(0);
```

```
}
```

POSIX

Pthread

"wait"

"1100"

1100

Win32 Thread

- Again Win32 thread is a specification, implementation varies among WinXP, Win7, etc.
- Win32 thread APIs are very similar to Pthread APIs
- Win32 threads are referred to as objects/handle
 - WaitForSingleObject
 - CloseHandle

```

#include <windows.h>
#include <stdio.h>
DWORD Sum; /* data is shared by the thread(s) */
/* the thread runs in this separate function */

DWORD WINAPI Summation(LPVOID Param)
{
    DWORD Upper = *(DWORD*)Param;
    for (DWORD i = 0; i <= Upper; i++)
        Sum += i;
    return 0;
}

int main(int argc, char *argv[])
{
    DWORD ThreadId;
    HANDLE ThreadHandle;
    int Param;
    /* perform some basic error checking */
    if (argc != 2) {
        fprintf(stderr, "An integer parameter is required\n");
        return -1;
    }
    Param = atoi(argv[1]);
    if (Param < 0) {
        fprintf(stderr, "An integer >= 0 is required\n");
        return -1;
    }

    // create the thread
    ThreadHandle = CreateThread(
        NULL, // default security attributes
        0, // default stack size
        Summation, // thread function
        &Param, // parameter to thread function
        0, // default creation flags
        &ThreadId); // returns the thread identifier

    if (ThreadHandle != NULL) {
        // now wait for the thread to finish
        WaitForSingleObject(ThreadHandle, INFINITE);

        // close the thread handle
        CloseHandle(ThreadHandle);

        printf("sum = %d\n", Sum);
    }
}

```

OPERATING-SYSTEM THREAD SUPPORT

Windows XP Threads

- Implements the one-to-one mapping
- Each thread contains
 - A thread id
 - Register set
 - Separate user and kernel stacks
 - Private data storage area
- The register set, stacks, and private storage area are known as the context of the threads
- The primary data structures of a thread include:
 - ETHREAD (executive thread block)
 - KTHREAD (kernel thread block)
 - TEB (thread environment block)

Linux Threads

- Linux refers to them as tasks rather than threads
- Thread creation is done through `clone()` system call
- `clone()` allows a child task to share the address space of the parent task (process)
 - But the stacks are separate
 - Different from `vfork()`

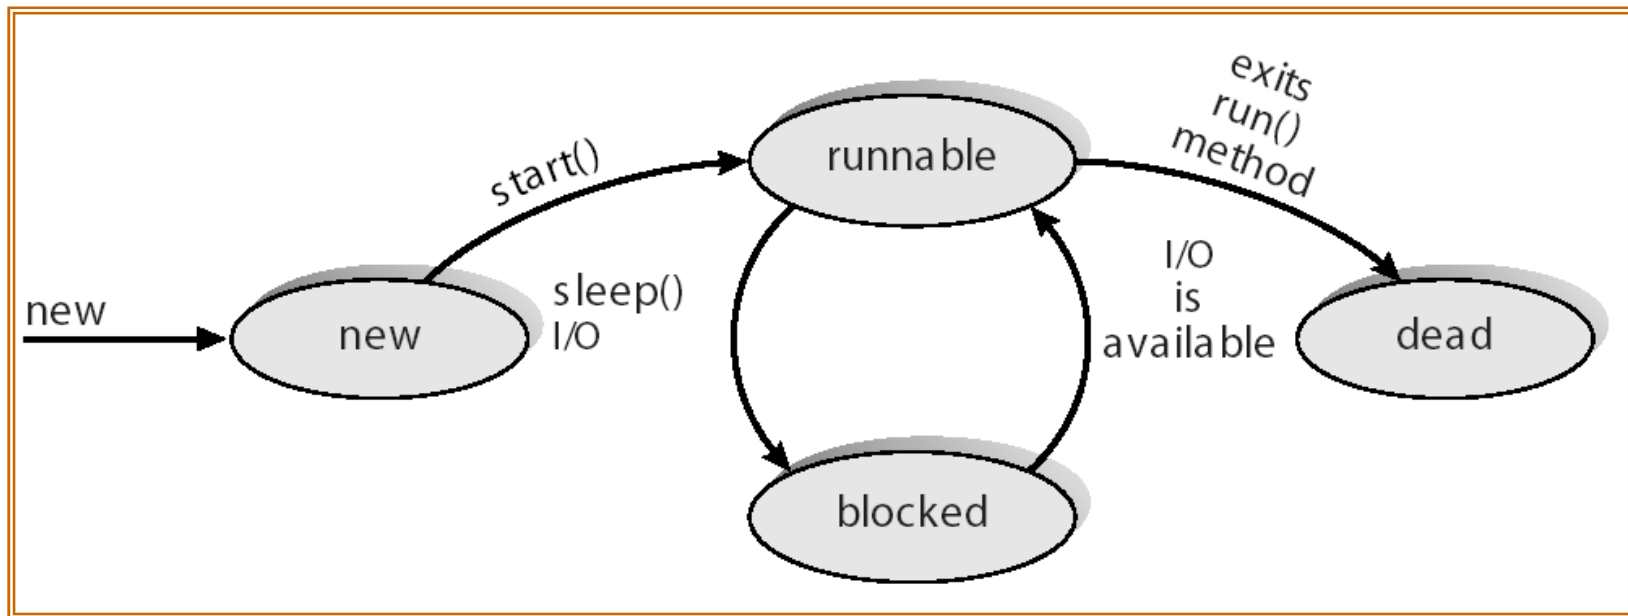
Java Threads

- Java threads are managed by the JVM
- Java threads may be created by:
 - Extending Thread class
 - Implementing the Runnable interface
- Threading model
 - Java green thread: many-to-one
 - Java (native) thread: many-to-many (old) or one-to-one (recent)

The JVM and the Host Operating System

The JVM is typically implemented on top of a host operating system (see Figure 2.20). This setup allows the JVM to hide the implementation details of the underlying operating system and to provide a consistent, abstract environment that allows Java programs to operate on any platform that supports a JVM. The specification for the JVM does not indicate how Java threads are to be mapped to the underlying operating system, instead leaving that decision to the particular implementation of the JVM. For example, the Windows XP operating system uses the one-to-one model; therefore, each Java thread for a JVM running on such a system maps to a kernel thread. On operating systems that use the many-to-many model (such as Tru64 UNIX), a Java thread is mapped according to the many-to-many model. Solaris initially implemented the JVM using the many-to-one model (the green threads library, mentioned earlier). Later releases of the JVM were implemented using the many-to-many model. Beginning with Solaris 9, Java threads were mapped using the one-to-one model. In addition, there may be a relationship between the Java thread library and the thread library on the host operating system. For example, implementations of a JVM for the Windows family of operating systems might use the Win32 API when creating Java threads; Linux, Solaris, and Mac OS X systems might use the Pthreads API.

Java Thread States



THREADING ISSUES

Semantics of fork() and exec()

- Does fork() duplicate only the calling thread or all threads?
 - Undefined, but in many UNIX variants the **entire process** (including all its threads) is **duplicated**
- How about exec()?
 - Again undefined, but in many UNIX variants the **entire process** (including all its threads) is **replaced**

Signal Handling

- Synchronous signal
 - Always delivered to the thread that causes the signal
 - E.g., access violation
- Asynchronous signal
 - Typically delivered to the first thread that dose not block the signal
 - E.g., process termination
- Varies from implementation to implementation

Thread Pools

- Create a number of threads in a pool where they await work
- Advantages:
 - Usually slightly faster to service a request with an **existing** thread than **create** a new thread
 - It is costly to repeatedly create and delete threads
 - Allows the number of threads in the application(s) to be bound to the size of the pool
 - Multiplexing tasks over threads, similar to the concept of the many-to-many model

End of Chapter 4